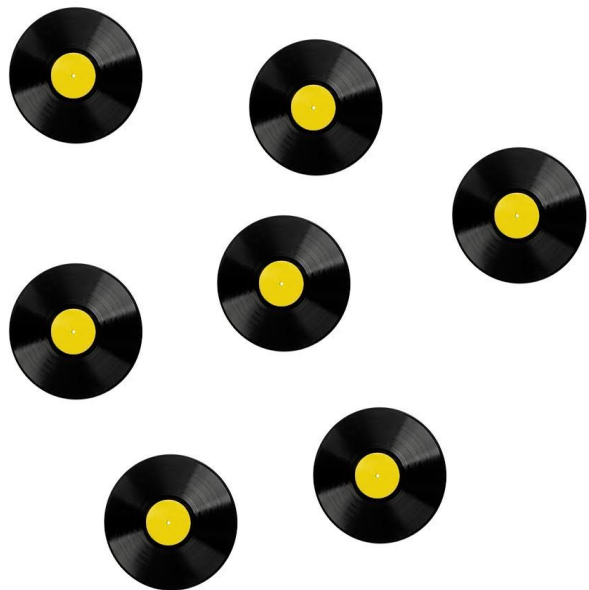


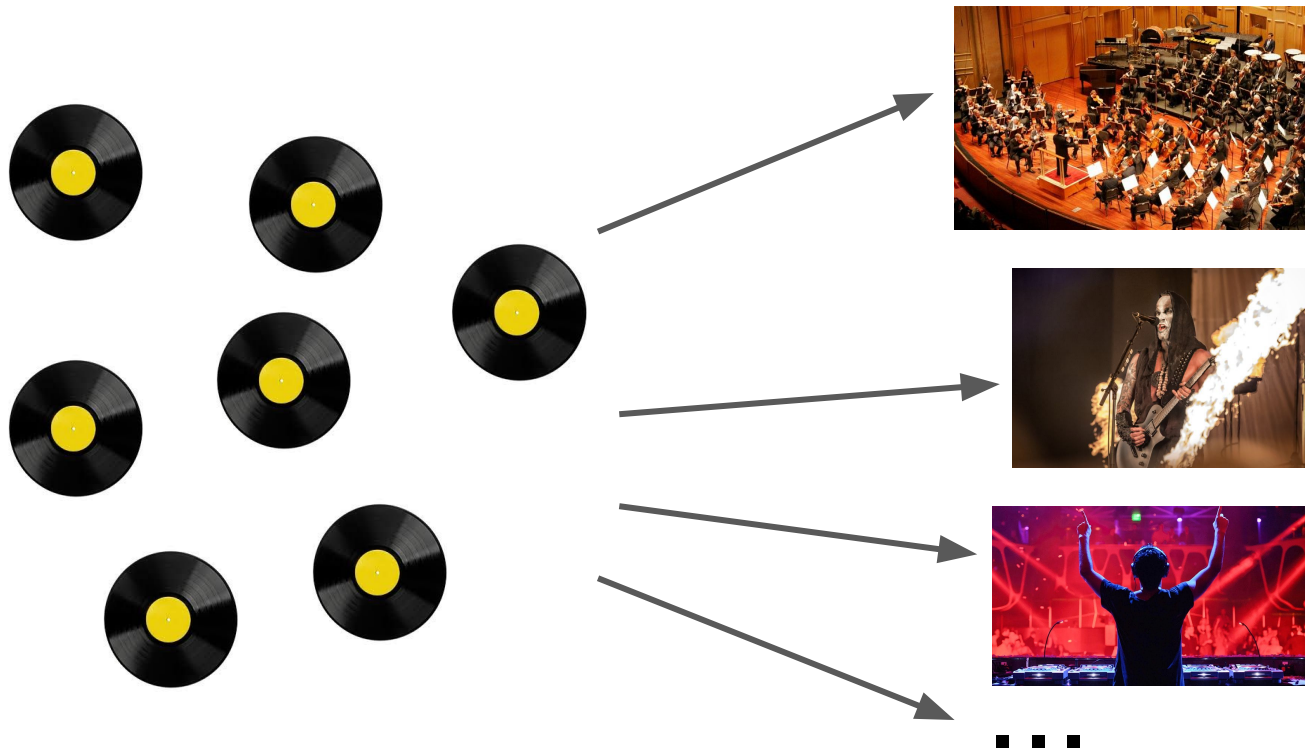
Implementing a neural network for music genre classification

Valerio Velardo

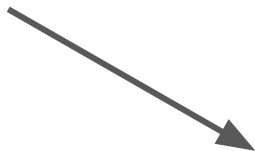
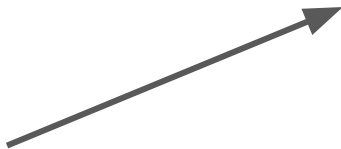
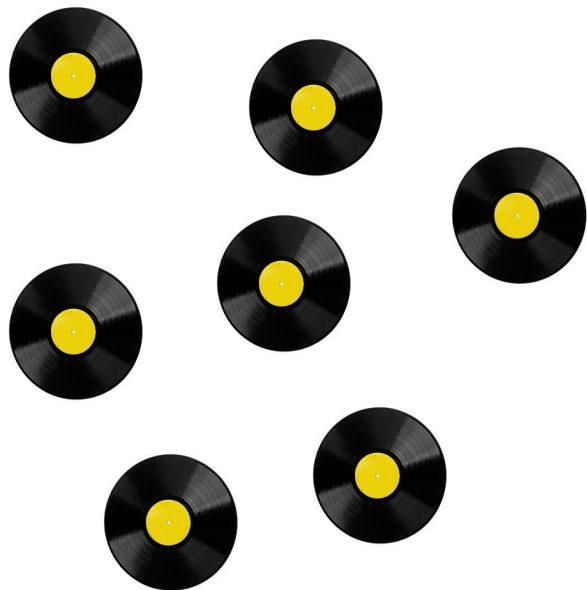
Multiclass classification



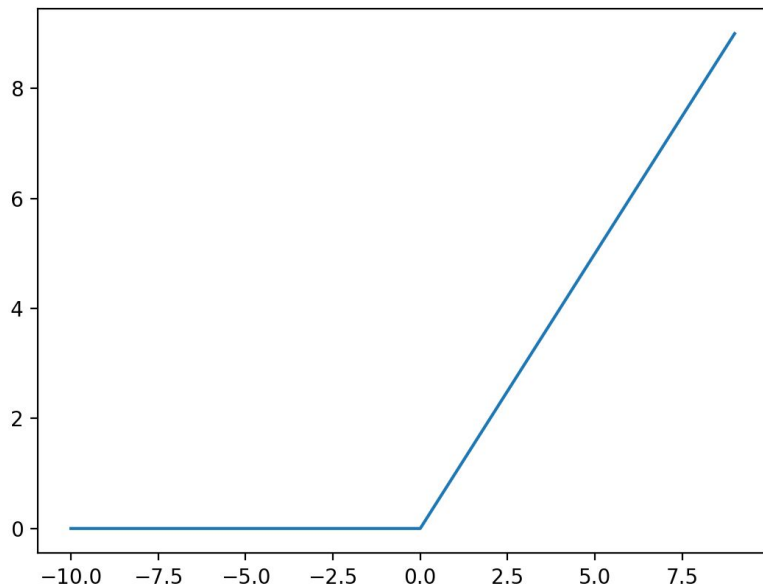
Multiclass classification



Binary classification



Rectified Linear Unit (ReLU)



$$ReLU(h) = \begin{cases} 0, & \text{if } h < 0 \\ h, & \text{if } h \geq 0 \end{cases}$$

Rectified Linear Unit (ReLU)

- Better convergence
- Reduced likelihood of *vanishing gradient*

Types of batching

- Stochastic
 - Calculate gradient on 1 sample
 - Quick, but inaccurate

Types of batching

- Stochastic
 - Calculate gradient on 1 sample
 - Quick, but inaccurate
- Full batch
 - Compute gradient on the whole training set
 - Slow, memory intensive, accurate

Types of batching

- Stochastic
 - Calculate gradient on 1 sample
 - Quick, but inaccurate
- Full batch
 - Compute gradient on the whole training set
 - Slow, memory intensive, accurate
- Mini-batch
 - Compute gradient on a subset of data set (16 - 128 samples)
 - Best of the 2 worlds